

# TravelTide Exploration

```
In [1]: from sqlalchemy import create_engine

# Create a connection using SQLAlchemy
DATABASE_URL = "postgresql+psycopg2://Test:bQNxVzJL4g6u@ep-noisy-flower-846766.us-east
engine = create_engine(DATABASE_URL)

In [2]: # Cohort filter definition
cohort_filter = """
WITH CohortUsers AS (
    SELECT user_id
    FROM sessions
    WHERE session_start > '2023-01-04'
    GROUP BY user_id
    HAVING COUNT(session_id) > 7
)
"""

In [3]: import pandas as pd
import matplotlib.pyplot as plt

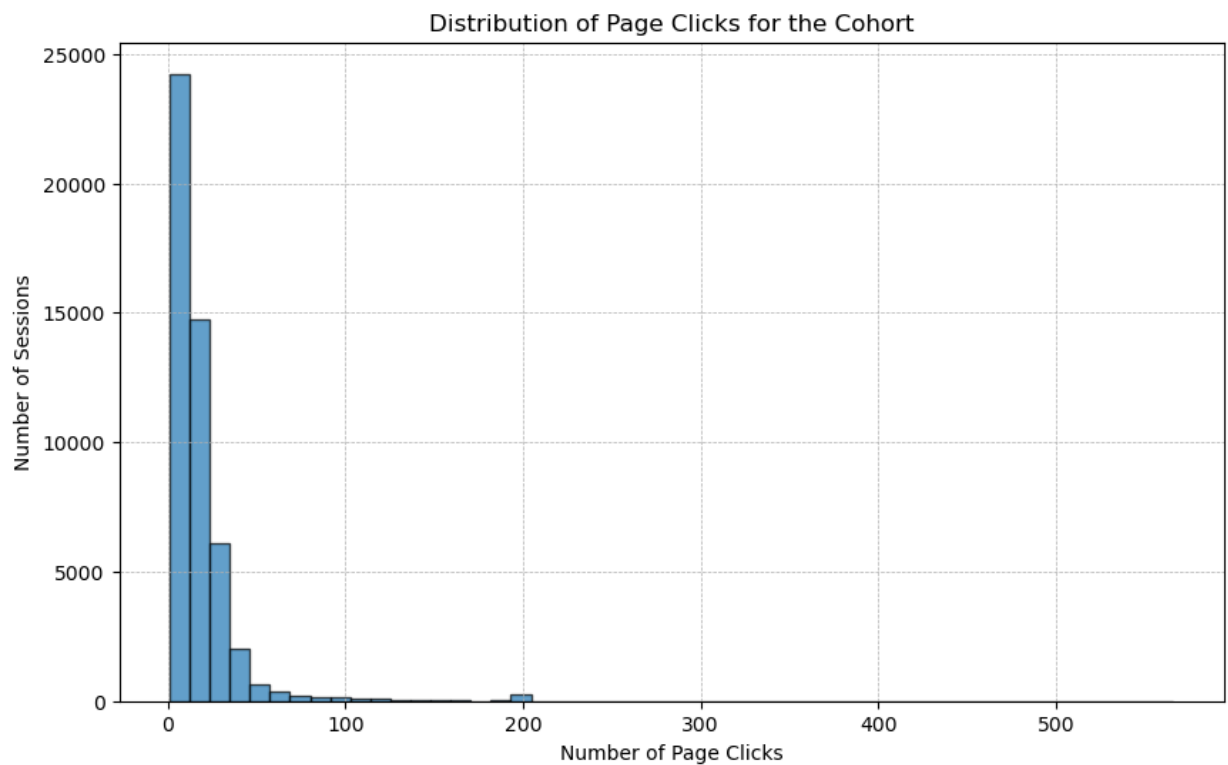
In [4]: # Page Clicks

query = f"""
{cohort_filter}
SELECT page_clicks
FROM sessions
WHERE session_start > '2023-01-04' AND user_id IN (SELECT user_id FROM CohortUsers)
"""

# Fetch the data
df_page_clicks = pd.read_sql(query, engine)

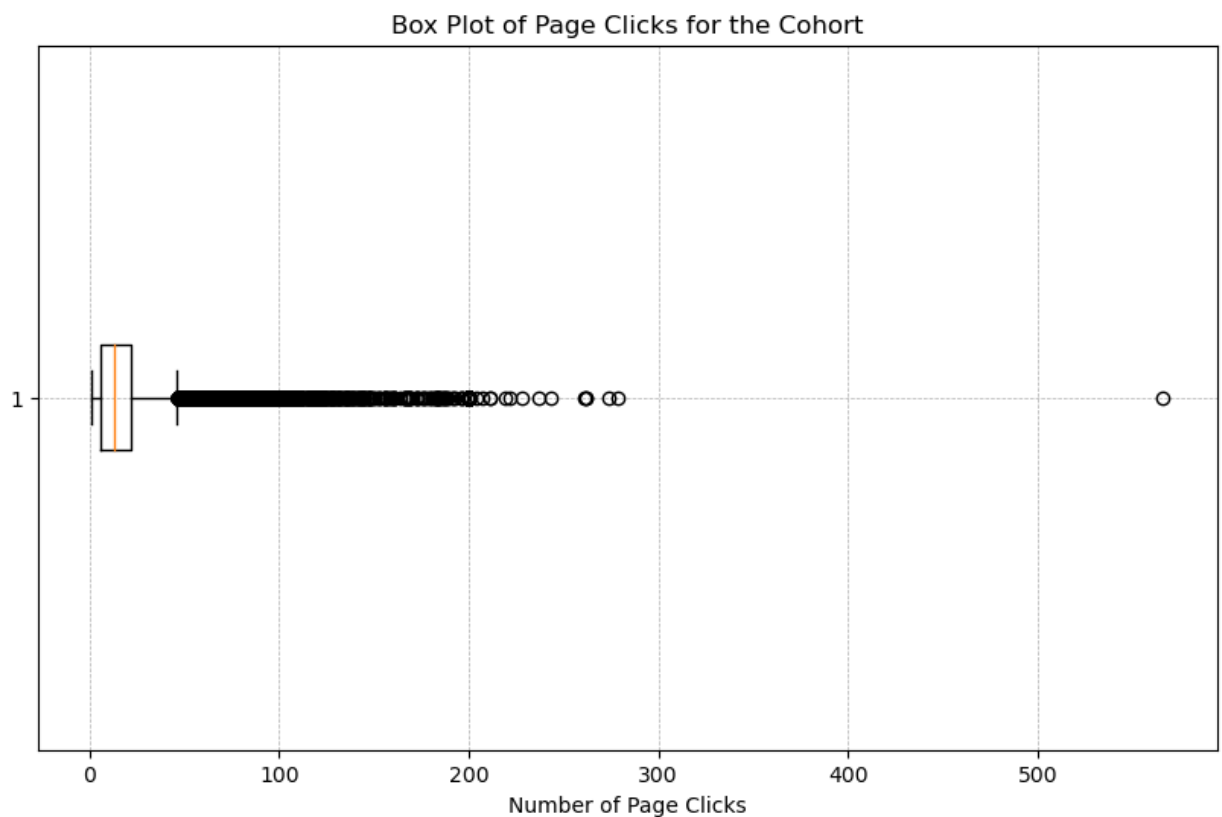
In [5]: # Histogram for Page clicks

plt.figure(figsize=(10, 6))
plt.hist(df_page_clicks['page_clicks'], bins=50, edgecolor='k', alpha=0.7)
plt.title('Distribution of Page Clicks for the Cohort')
plt.xlabel('Number of Page Clicks')
plt.ylabel('Number of Sessions')
plt.grid(True, which='both', linestyle='--', linewidth=0.5)
plt.show()
```



```
In [6]: # Box Plot for Page Clicks

plt.figure(figsize=(10, 6))
plt.boxplot(df_page_clicks['page_clicks'], vert=False)
plt.title('Box Plot of Page Clicks for the Cohort')
plt.xlabel('Number of Page Clicks')
plt.grid(True, which='both', linestyle='--', linewidth=0.5)
plt.show()
```



```
In [7]: # Number of Outliers in Page clicks

# Calculate Q1, Q3 and IQR
Q1 = df_page_clicks['page_clicks'].quantile(0.25)
Q3 = df_page_clicks['page_clicks'].quantile(0.75)
IQR = Q3 - Q1

# Determine boundaries
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filter the dataframe for outliers and count them
outliers = df_page_clicks[(df_page_clicks['page_clicks'] < lower_bound) | (df_page_clicks['page_clicks'] > upper_bound)]
num_outliers = outliers.shape[0]

# Total number of data points
total_data_points = len(df_page_clicks)

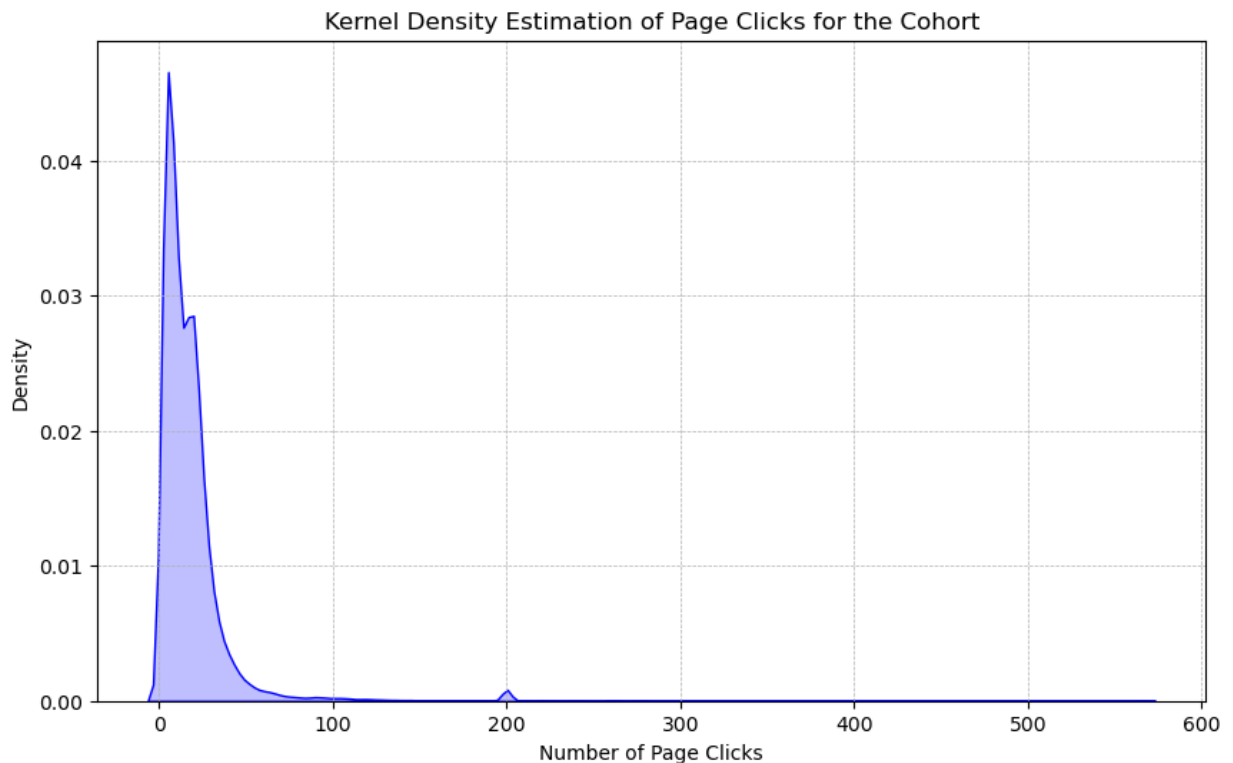
# Calculate the percentage
percentage_outliers = (num_outliers / total_data_points) * 100

print(f"Number of outliers: {num_outliers}")
print(f"Percentage of outliers: {percentage_outliers}")
```

Number of outliers: 2103  
Percentage of outliers: 4.273434801162342

```
In [8]: # KDE Plot for Page Clicks
import seaborn as sns

plt.figure(figsize=(10, 6))
sns.kdeplot(df_page_clicks['page_clicks'], fill=True, color='blue')
plt.title('Kernel Density Estimation of Page Clicks for the Cohort')
plt.xlabel('Number of Page Clicks')
plt.ylabel('Density')
plt.grid(True, which='both', linestyle='--', linewidth=0.5)
plt.show()
```



```
In [9]: # Skewness and Kurtosis
from scipy.stats import skew, kurtosis

skewness = skew(df_page_clicks['page_clicks'])
kurt = kurtosis(df_page_clicks['page_clicks'])
print(f"Skewness: {skewness}")
print(f"Kurtosis: {kurt}")
```

```
Skewness: 5.3902993496060505
Kurtosis: 44.364076329644234
```

```
In [10]: # Descriptive Statistics
print(df_page_clicks['page_clicks'].describe())
```


```
count    49211.000000
mean      17.588791
std       21.495987
min        1.000000
25%        6.000000
50%       13.000000
75%       22.000000
max       566.000000
Name: page_clicks, dtype: float64
```

```
In [11]: # If skewness is significantly different from 0, consider transformations
```

```
import numpy as np

if abs(skewness) > 1:
    # Log transformation (adding 1 to handle zero values)
    df_page_clicks['log_page_clicks'] = df_page_clicks['page_clicks'].apply(lambda x:
    plt.figure(figsize=(10, 6))
    sns.kdeplot(df_page_clicks['log_page_clicks'], fill=True, color='green')
    plt.title('Kernel Density Estimation of Log-transformed Page Clicks for the Cohort')
    plt.xlabel('Log-transformed Number of Page Clicks')
```

```
plt.ylabel('Density')
plt.grid(True, which='both', linestyle='--', linewidth=0.5)
plt.show()
```



```
In [12]: # Let's get the Session Duration now

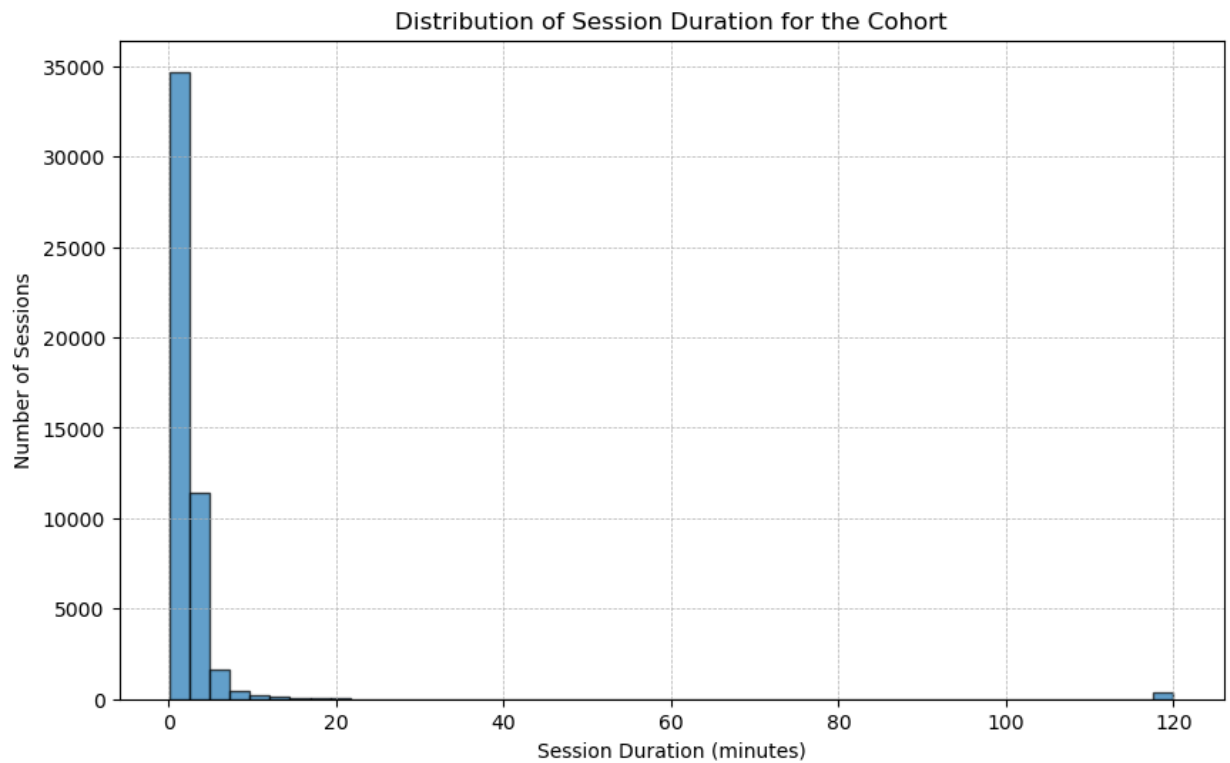
query_duration = f"""
{cohort_filter}
SELECT session_start, session_end
FROM sessions
WHERE session_start > '2023-01-04' AND user_id IN (SELECT user_id FROM CohortUsers)
"""

# Fetch the data
df_session_times = pd.read_sql(query_duration, engine)

# Calculate session duration
df_session_times['session_duration'] = (df_session_times['session_end'] - df_session_t
```

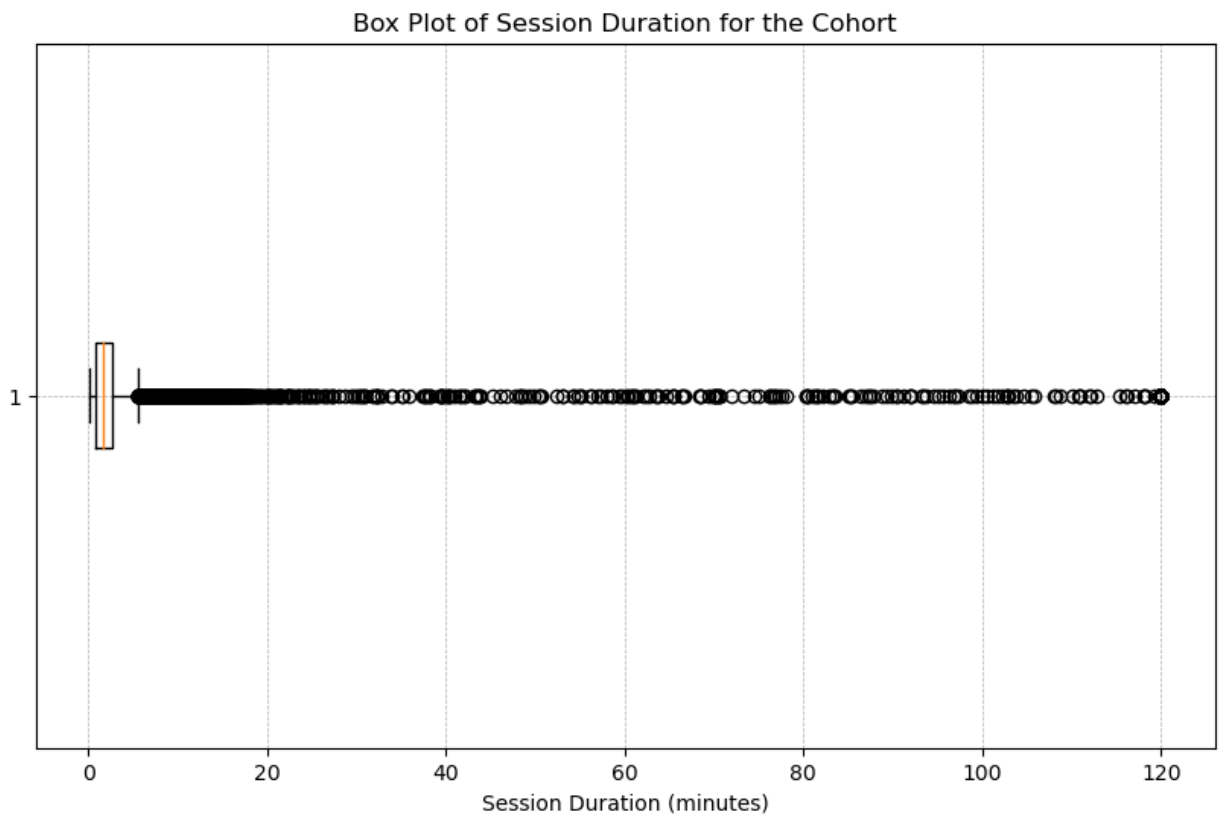
```
In [13]: # Histogram for duration

plt.figure(figsize=(10, 6))
plt.hist(df_session_times['session_duration'], bins=50, edgecolor='k', alpha=0.7)
plt.title('Distribution of Session Duration for the Cohort')
plt.xlabel('Session Duration (minutes)')
plt.ylabel('Number of Sessions')
plt.grid(True, which='both', linestyle='--', linewidth=0.5)
plt.show()
```




```
In [14]: # Box Plot for Session Duration

plt.figure(figsize=(10, 6))
plt.boxplot(df_session_times['session_duration'], vert=False)
plt.title('Box Plot of Session Duration for the Cohort')
plt.xlabel('Session Duration (minutes)')
plt.grid(True, which='both', linestyle='--', linewidth=0.5)
plt.show()
```



In [15]: *# KDE Plot for Session Duration*

```
plt.figure(figsize=(10, 6))
sns.kdeplot(df_session_times['session_duration'], fill=True, color='blue')
plt.title('KDE Plot of Session Duration for the Cohort')
plt.xlabel('Session Duration (minutes)')
plt.ylabel('Density')
plt.grid(True, which='both', linestyle='--', linewidth=0.5)
plt.show()
```



In [16]: *# Number of Outliers in Session Duration*

```
# Calculate Q1, Q3 and IQR
Q1_duration = df_session_times['session_duration'].quantile(0.25)
Q3_duration = df_session_times['session_duration'].quantile(0.75)
IQR_duration = Q3_duration - Q1_duration

# Determine boundaries
lower_bound_duration = Q1_duration - 1.5 * IQR_duration
upper_bound_duration = Q3_duration + 1.5 * IQR_duration

# Filter the dataframe for outliers and count them
outliers_duration = df_session_times[(df_session_times['session_duration'] < lower_bound_duration) | (df_session_times['session_duration'] > upper_bound_duration)]
num_outliers_duration = outliers_duration.shape[0]

# Total number of data points
total_data_points_duration = len(df_session_times)

# Calculate the percentage
percentage_outliers_duration = (num_outliers_duration / total_data_points_duration) * 100

print(f'Number of outliers in session duration: {num_outliers_duration}')
print(f'Percentage of outliers in session duration: {percentage_outliers_duration}')
```

Number of outliers in session duration: 2332

Percentage of outliers in session duration: 4.738777915506696

## Let's spend some time looking at users

```
In [17]: # SQL query to find the number of unique users in the cohort
query_unique_users_in_cohort = f"""
{cohort_filter}
SELECT COUNT(DISTINCT user_id) as unique_users
FROM CohortUsers
"""

# Fetch the data
df_unique_users_in_cohort = pd.read_sql(query_unique_users_in_cohort, engine)

# Display the result
print(f"Number of unique users in the cohort: {df_unique_users_in_cohort['unique_users']}")
```

Number of unique users in the cohort: 5998

```
In [18]: # SQL query to find the number of unique users in the cohort who made a booking
query_users_with_bookings = f"""
{cohort_filter}
SELECT COUNT(DISTINCT s.user_id) as users_with_bookings
FROM sessions s
INNER JOIN CohortUsers c ON s.user_id = c.user_id
WHERE s.flight_booked = True OR s.hotel_booked = True
"""

# Fetch the data
df_users_with_bookings = pd.read_sql(query_users_with_bookings, engine)

# Display the result
print(f"Number of unique users in the cohort who made a booking: {df_users_with_bookings['users_with_bookings']}")
```

Number of unique users in the cohort who made a booking: 5566

```
In [19]: # SQL query to find users who frequently book hotels
query_frequent_hotel_bookers = f"""
{cohort_filter}
SELECT s.user_id, COUNT(*) as num_hotel_bookings
FROM sessions s
INNER JOIN CohortUsers c ON s.user_id = c.user_id
WHERE s.hotel_booked = True
GROUP BY s.user_id
ORDER BY num_hotel_bookings DESC
"""

df_frequent_hotel_bookers = pd.read_sql(query_frequent_hotel_bookers, engine)

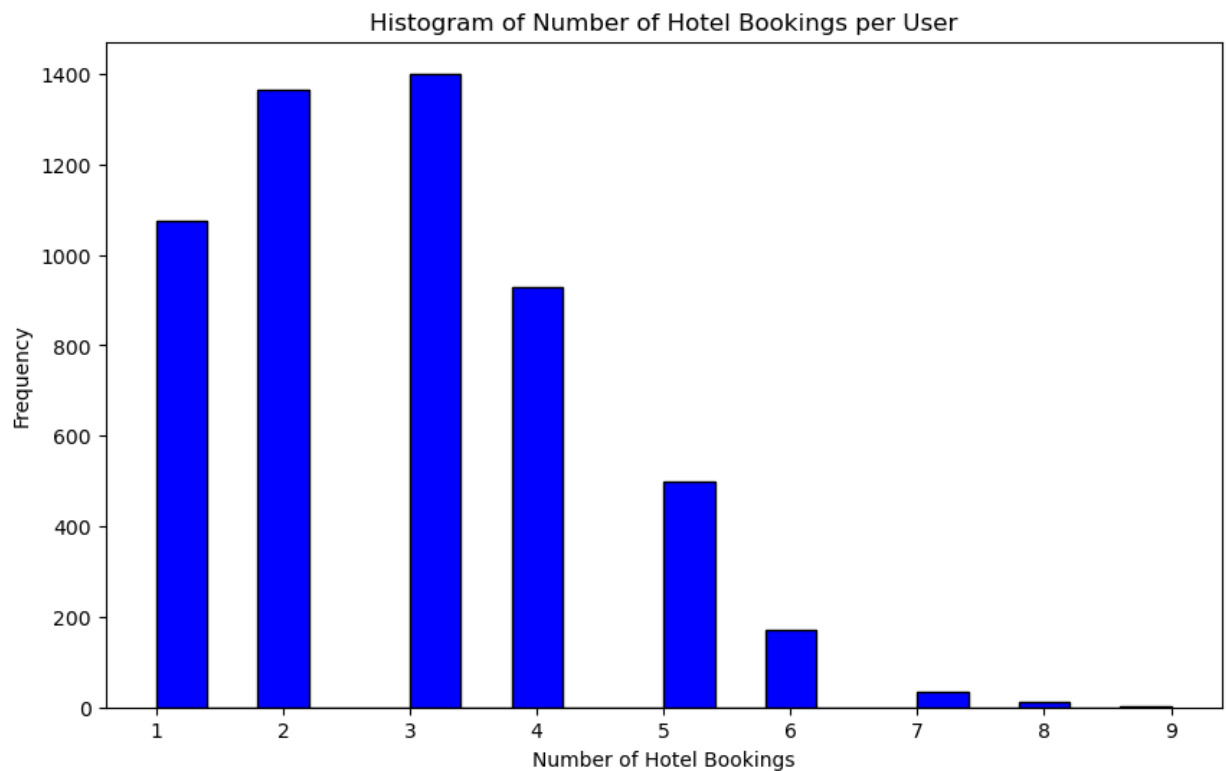
# Descriptive Statistics
print(df_frequent_hotel_bookers['num_hotel_bookings'].describe())


# Histogram
plt.figure(figsize=(10, 6))
plt.hist(df_frequent_hotel_bookers['num_hotel_bookings'], bins=20, color='blue', edgecolor='black')
plt.xlabel('Number of Hotel Bookings')
plt.ylabel('Frequency')
plt.title('Histogram of Number of Hotel Bookings per User')
plt.show()
```



```
# Box Plot
plt.figure(figsize=(10, 6))
plt.boxplot(df_frequent_hotel_bookers['num_hotel_bookings'])
plt.ylabel('Number of Hotel Bookings')
plt.title('Box Plot of Number of Hotel Bookings per User')
plt.show()
```

```
count    5486.000000
mean       2.842144
std        1.411458
min         1.000000
25%         2.000000
50%         3.000000
75%         4.000000
max         9.000000
Name: num_hotel_bookings, dtype: float64
```





```
In [20]: # SQL query to find users who frequently check bags, but now with average checked bags
query_frequent_bag_checkers = f"""
{cohort_filter},
FilteredSessions AS (
    SELECT s.user_id, s.trip_id
    FROM sessions s
    INNER JOIN CohortUsers c ON s.user_id = c.user_id
),
BagCounts AS (
    SELECT s.user_id, f.trip_id, COUNT(*) as num_checked_bags
    FROM flights f
    INNER JOIN FilteredSessions s ON f.trip_id = s.trip_id
    GROUP BY s.user_id, f.trip_id
)
SELECT user_id, AVG(num_checked_bags) as avg_checked_bags_per_trip
FROM BagCounts
GROUP BY user_id
ORDER BY avg_checked_bags_per_trip DESC;
"""

df_frequent_bag_checkers = pd.read_sql(query_frequent_bag_checkers, engine)
print("sql done")
# Descriptive Statistics
print(df_frequent_bag_checkers['avg_checked_bags_per_trip'].describe())

# Histogram
plt.figure(figsize=(10, 6))
plt.hist(df_frequent_bag_checkers['avg_checked_bags_per_trip'], bins=20, color='green')
plt.xlabel('Average Number of Checked Bags per Trip')
plt.ylabel('Frequency')
plt.title('Histogram of Average Number of Checked Bags per Trip per User')
plt.show()

# Box Plot
plt.figure(figsize=(10, 6))
```

```
plt.boxplot(df_frequent_bag_checkers['avg_checked_bags_per_trip'])  
plt.ylabel('Average Number of Checked Bags per Trip')  
plt.title('Box Plot of Average Number of Checked Bags per Trip per User')  
plt.show()
```

sql done

count 5206.000000

mean 1.053488

std 0.179770

min 1.000000

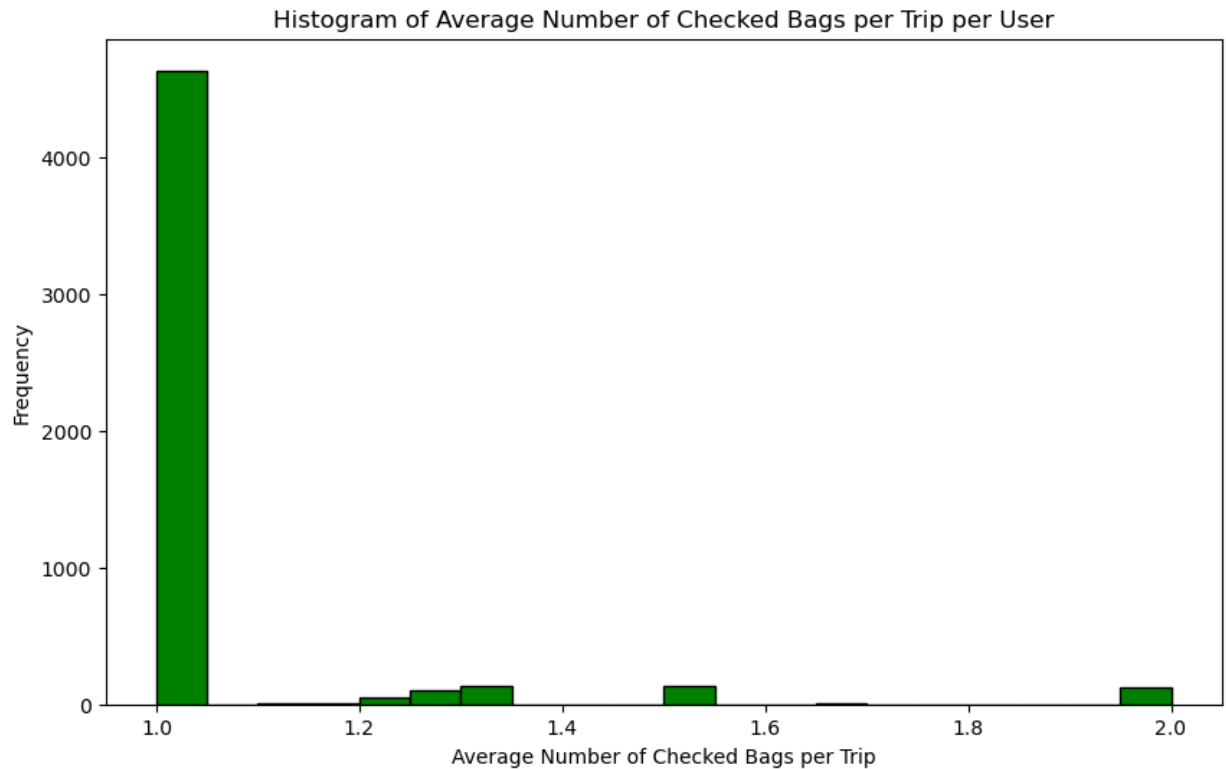
25% 1.000000

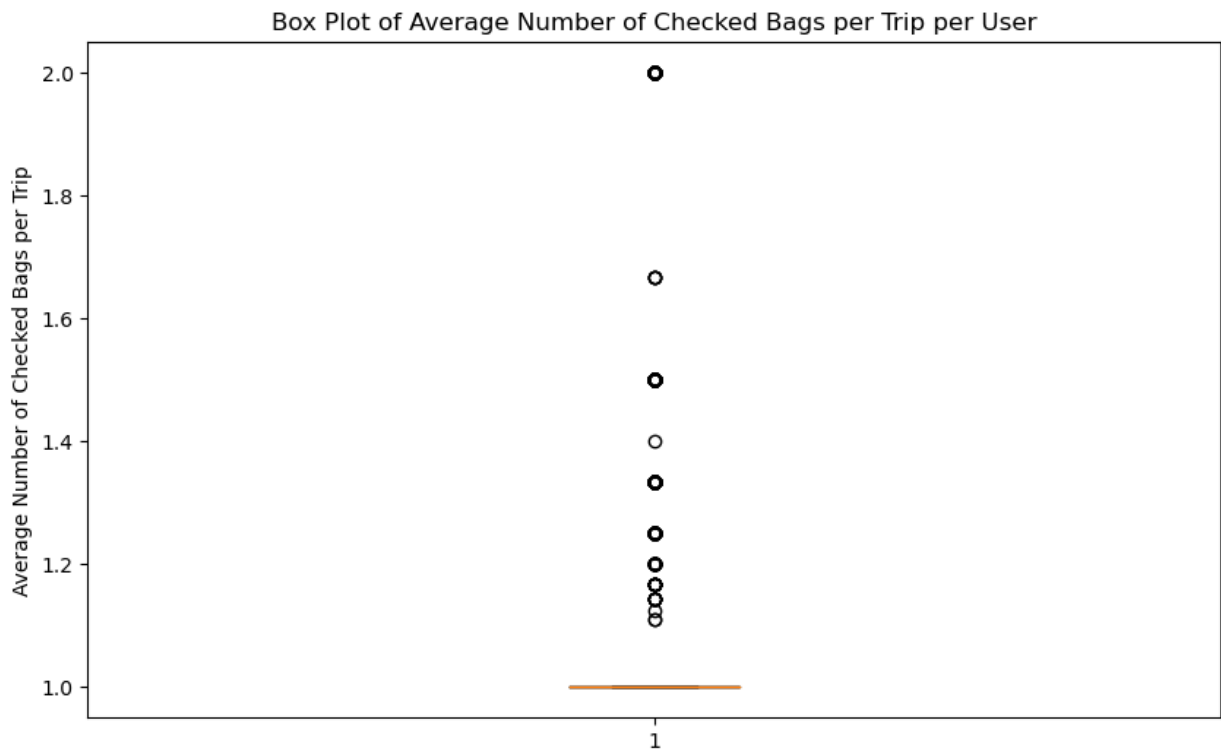
50% 1.000000

75% 1.000000

max 2.000000

Name: avg\_checked\_bags\_per\_trip, dtype: float64





Almost every user checks at least one bag, this is not an insight that will contribute to our segmentation. Let's check other aspects

## Let's find the high frequency travellers

```
In [21]: # SQL query to find high frequency travelers in the cohort
query_high_frequency_travelers = f"""
{cohort_filter}
SELECT s.user_id, COUNT(DISTINCT s.trip_id) as num_trips
FROM sessions s
INNER JOIN CohortUsers c ON s.user_id = c.user_id
GROUP BY s.user_id
ORDER BY num_trips DESC
"""

# Execute the query and store the result in a DataFrame
df_high_frequency_travelers = pd.read_sql(query_high_frequency_travelers, engine)

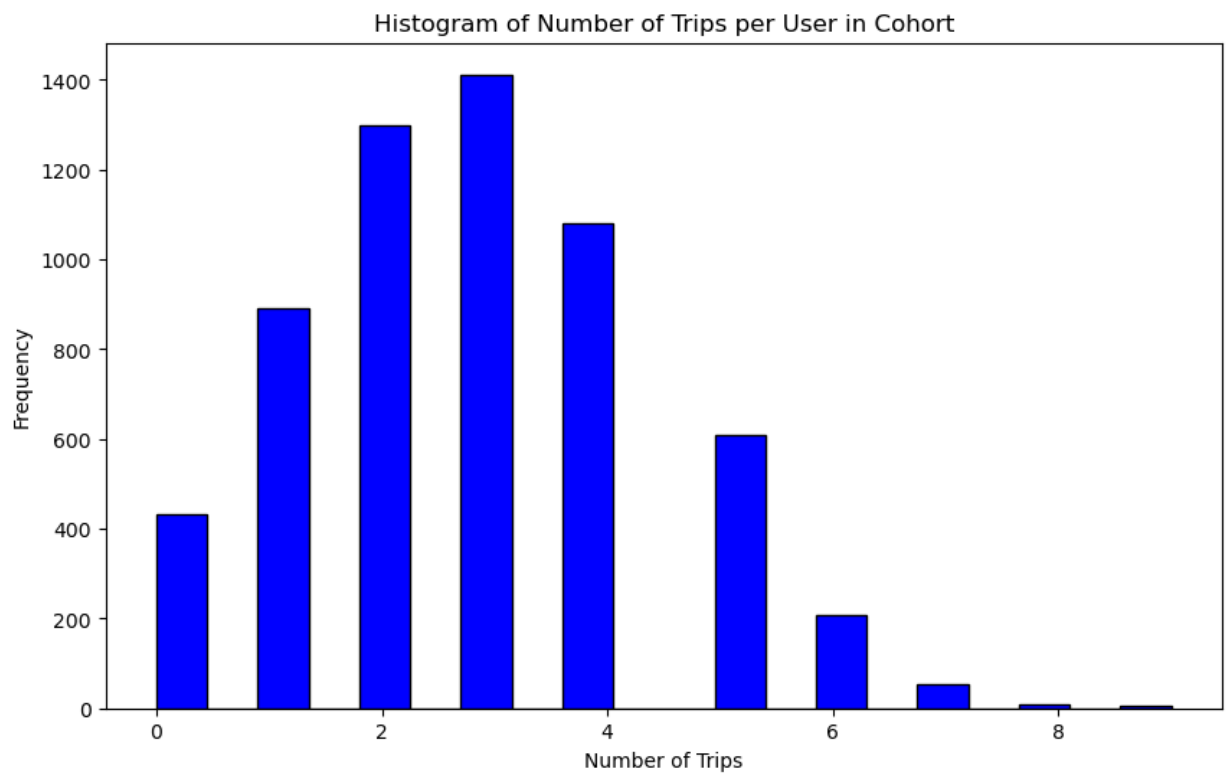
# Display descriptive statistics
print(df_high_frequency_travelers['num_trips'].describe())


# Plotting
# Histogram
plt.figure(figsize=(10, 6))
plt.hist(df_high_frequency_travelers['num_trips'], bins=20, color='blue', edgecolor='k')
plt.xlabel('Number of Trips')
plt.ylabel('Frequency')
plt.title('Histogram of Number of Trips per User in Cohort')
plt.show()

# Box Plot
plt.figure(figsize=(10, 6))
plt.boxplot(df_high_frequency_travelers['num_trips'])
```

```
plt.ylabel('Number of Trips')  
plt.title('Box Plot of Number of Trips per User in Cohort')  
plt.show()
```

```
count    5998.000000  
mean      2.805435  
std       1.589629  
min       0.000000  
25%       2.000000  
50%       3.000000  
75%       4.000000  
max       9.000000  
Name: num_trips, dtype: float64
```





## Long Distance Flights

Vincenty Function to replace the Haversine function - Vincenty is much more accurate

[https://en.wikipedia.org/wiki/Vincenty%27s\\_formulae](https://en.wikipedia.org/wiki/Vincenty%27s_formulae)

```
In [22]: from math import radians, sin, cos, sqrt, atan2

def vincenty_distance(lat1, lon1, lat2, lon2):
    """
    Calculate the great-circle distance between two points
    on the Earth surface given their latitude and longitude
    in decimal degrees.
    """

    # WGS-84 ellipsoid parameters
    a = 6378137.0 # semi-major axis in meters
    f = 1 / 298.257223563 # flattening
    b = (1 - f) * a # semi-minor axis

    # convert decimal degrees to radians
    lat1, lon1, lat2, lon2 = map(radians, [lat1, lon1, lat2, lon2])

    # calculations
    U1 = atan((1 - f) * tan(lat1))
    U2 = atan((1 - f) * tan(lat2))
    sinU1 = sin(U1)
    cosU1 = cos(U1)
    sinU2 = sin(U2)
    cosU2 = cos(U2)
```

```

lon_diff = lon2 - lon1
Lambda = lon_diff # initial approximation for Lambda
sinLambda = sin(Lambda)
cosLambda = cos(Lambda)

# iterate until change is insignificant
for _ in range(1000):
    sinSigma = sqrt((cosU2 * sin(Lambda)) ** 2 + (cosU1 * sinU2 - sinU1 * cosU2 *
    cosSigma = sinU1 * sinU2 + cosU1 * cosU2 * cos(Lambda)
    sigma = atan2(sinSigma, cosSigma)
    sinAlpha = cosU1 * cosU2 * sin(Lambda) / sinSigma
    cos2Alpha = 1 - sinAlpha ** 2
    cos2SigmaM = cosSigma - 2 * sinU1 * sinU2 / cos2Alpha
    C = f / 16 * cos2Alpha * (4 + f * (4 - 3 * cos2Alpha))
    Lambda_prev = Lambda
    Lambda = lon_diff + (1 - C) * f * sinAlpha * (sigma + C * sinSigma * (cos2SigmaM
    # break if change in lambda is insignificant
    if abs(Lambda - Lambda_prev) < 1e-12:
        break

# final calculations
u2 = cos2Alpha * (a ** 2 - b ** 2) / (b ** 2)
A = 1 + u2 / 16384 * (4096 + u2 * (-768 + u2 * (320 - 175 * u2)))
B = u2 / 1024 * (256 + u2 * (-128 + u2 * (74 - 47 * u2)))
deltaSigma = B * sinSigma * (cos2SigmaM + B / 4 * (cosSigma * (-1 + 2 * cos2SigmaM

# distance in meters
s = b * A * (sigma - deltaSigma)

return s

```

## Let's explore Flight Distance

```

In [23]: # SQL query to get flight data
flight_query = f"""
{cohort_filter}
SELECT
    f.origin_airport,
    f.destination_airport,
    u.home_airport_lat AS origin_airport_lat,
    u.home_airport_lon AS origin_airport_lon,
    f.destination_airport_lat,
    f.destination_airport_lon
FROM flights f
JOIN sessions s ON f.trip_id = s.trip_id
JOIN CohortUsers c ON s.user_id = c.user_id
JOIN users u ON s.user_id = u.user_id
"""

# Execute the query and load the data into a DataFrame
df_flights = pd.read_sql(flight_query, engine)

```

```

In [24]: !pip install geopy

```

Requirement already satisfied: geopy in c:\users\bhaze\anaconda3\lib\site-packages (2.4.0)

Requirement already satisfied: geographiclib<3,>=1.52 in c:\users\bhaze\anaconda3\lib\site-packages (from geopy) (2.0)

In [25]: `from geopy.distance import geodesic`

```
# Initialize an empty List to store the distances
distances = []
```

```
# Loop through the DataFrame to calculate distances
```

```
for index, row in df_flights.iterrows():
    origin = (row['origin_airport_lat'], row['origin_airport_lon'])
    destination = (row['destination_airport_lat'], row['destination_airport_lon'])
    distance = geodesic(origin, destination).kilometers
    distances.append(distance)
```

```
# Add the distances to the DataFrame
```

```
df_flights['trip_distance_km'] = distances
```

```
# Display the updated DataFrame
```

```
print(df_flights.head())
```

	origin_airport	destination_airport	origin_airport_lat	origin_airport_lon	\
0	IND	YMX	39.717	-86.294	
1	LAX	YVR	33.942	-118.408	
2	LGA	YZD	40.777	-73.872	
3	IAD	JFK	38.944	-77.456	
4	LUK	DEN	39.103	-84.419	

	destination_airport_lat	destination_airport_lon	trip_distance_km
0	45.517	-73.417	1235.221296
1	49.195	-123.182	1739.332492
2	43.862	-79.370	567.979189
3	40.640	-73.779	366.896751
4	39.858	-104.667	1740.377783

In [26]: `# Display descriptive statistics`

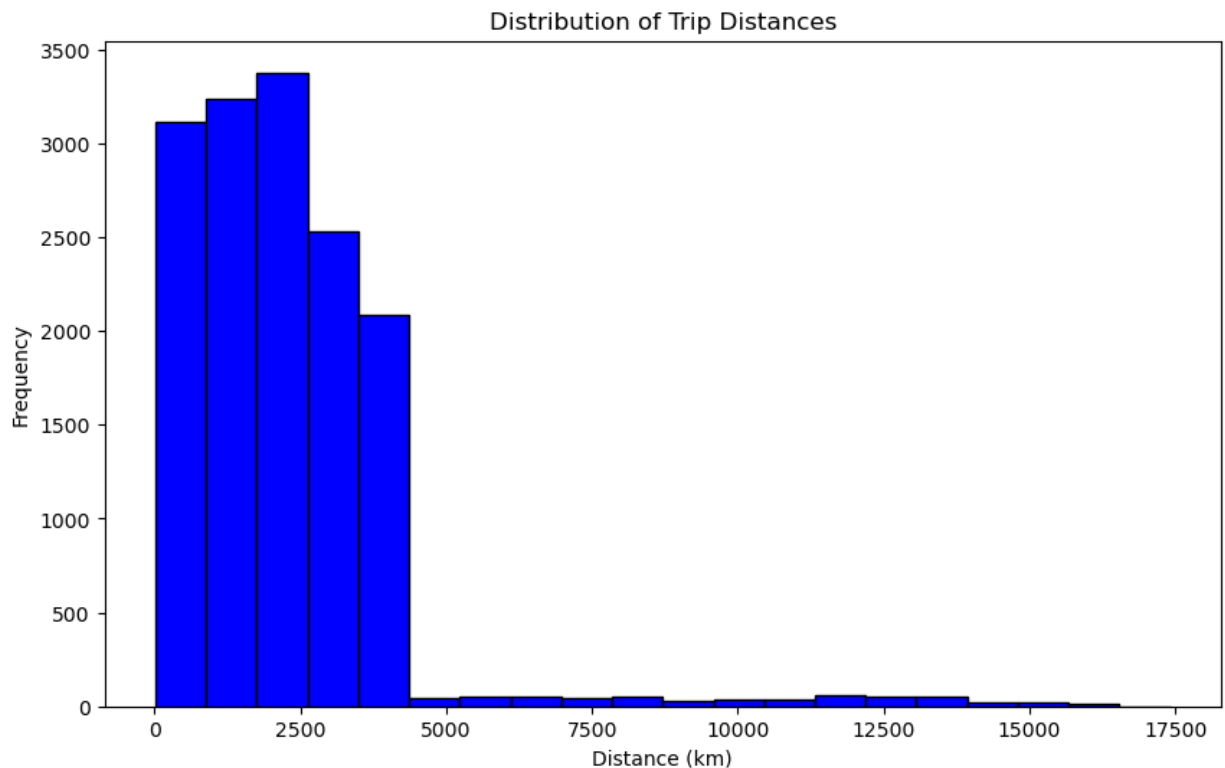
```
print(df_flights['trip_distance_km'].describe())
```

```
count    14919.000000
mean      2331.347257
std       1988.553779
min        17.685030
25%       1069.915667
50%       2064.546063
75%       3132.946871
max       17407.930322
Name: trip_distance_km, dtype: float64
```

In [27]: `# Create a histogram for trip distances`

```
plt.figure(figsize=(10, 6))
plt.hist(df_flights['trip_distance_km'], bins=20, color='blue', edgecolor='black')
plt.title('Distribution of Trip Distances')
plt.xlabel('Distance (km)')
plt.ylabel('Frequency')
plt.show()
```





We're going to go off track and start looking at data in a different way. I'm really going to need to clean up this notebook!

## Free Hotel Meal segemntation

- Budget-Conscious Travelers
- International Travelers
- Couples on Weekend Getaways

## Budget-Conscious Travelers

We are going to look at a few aspects of the data. Some can be reused in other explorations:

1. Average Base Fare for Flights
2. Average Hotel Per Room Cost
3. Usage of Discounts

```
In [28]: # Get the average base fares
avg_base_fare_query = f"""
{cohort_filter}
SELECT AVG(f.base_fare_usd) AS avg_base_fare, s.user_id
FROM flights f
JOIN sessions s ON f.trip_id = s.trip_id
JOIN CohortUsers c ON s.user_id = c.user_id
GROUP BY s.user_id;
```

```
"""
df_avg_base_fare = pd.read_sql(avg_base_fare_query, engine)

# Get the number of bookings with a discount
discount_booking_query = f"""
{cohort_filter}
SELECT COUNT(*) AS num_discounts, s.user_id
FROM sessions s
JOIN CohortUsers c ON s.user_id = c.user_id
WHERE CAST(s.flight_discount AS INTEGER) = 1 OR CAST(s.hotel_discount AS INTEGER) = 1
GROUP BY s.user_id;
"""

df_discount_booking = pd.read_sql(discount_booking_query, engine)

# Get the Average Price for Hotel Stays
average_price_hotel_stays_query = f"""
{cohort_filter}
SELECT AVG(h.hotel_per_room_usd) AS avg_hotel_price, s.user_id
FROM hotels h
JOIN sessions s ON h.trip_id = s.trip_id
JOIN CohortUsers c ON s.user_id = c.user_id
GROUP BY s.user_id;
"""

df_average_hotel_stays = pd.read_sql(average_price_hotel_stays_query, engine)
```

In [29]: *# Descriptive Statistics first*

```
print(df_avg_base_fare['avg_base_fare'].describe())
print(df_discount_booking['num_discounts'].describe())
print(df_average_hotel_stays['avg_hotel_price'].describe())
```

```

count      5206.000000
mean       538.526171
std        685.388137
min         5.350000
25%        284.835000
50%        392.895000
75%        564.665375
max        14280.380000
Name: avg_base_fare, dtype: float64
count      5618.000000
mean        2.533464
std         1.204718
min         1.000000
25%         2.000000
50%         2.000000
75%         3.000000
max         7.000000
Name: num_discounts, dtype: float64
count      5435.000000
mean       178.336311
std        84.305174
min        24.000000
25%       124.500000
50%       163.400000
75%       212.000000
max       1063.000000
Name: avg_hotel_price, dtype: float64

```

What the heck is with a \$5.35 flight??? This probably won't add anything to the analysis, but I need to find out

```

In [30]: # SQL query to investigate the $5.35 flight with Latitude and Longitude
investigate_cheap_flight_query = """
SELECT f.base_fare_usd, u.home_airport_lat AS origin_lat, u.home_airport_lon AS origin_lon,
       f.destination_airport_lat AS dest_lat, f.destination_airport_lon AS dest_lon, f
FROM flights f
JOIN sessions s ON f.trip_id = s.trip_id
JOIN users u ON s.user_id = u.user_id
WHERE f.base_fare_usd = 5.35
ORDER BY s.user_id;
"""

# Execute the query and store the result in a DataFrame
df_cheap_flight = pd.read_sql(investigate_cheap_flight_query, engine)

# Calculate the distance using vincenty_distance
#df_cheap_flight['calculated_distance'] = df_cheap_flight.apply(
#    lambda row: vincenty_distance((row['origin_lat'], row['origin_lon']), (row['dest_
#    #)

# Display the DataFrame with calculated distance
print(df_cheap_flight)

```

	base_fare_usd	origin_lat	origin_lon	dest_lat	dest_lon	\
0	5.35	33.818	-118.151	33.942	-118.408	
1	5.35	33.676	-117.868	33.942	-118.408	
2	5.35	33.818	-118.151	33.942	-118.408	
3	5.35	33.818	-118.151	33.942	-118.408	
4	5.35	33.818	-118.151	33.942	-118.408	
5	5.35	33.818	-118.151	33.942	-118.408	
6	5.35	40.692	-74.169	40.640	-73.779	
7	5.35	42.422	-87.868	42.947	-87.896	

	departure_time	return_flight_booked	return_time	seats	\
0	2022-03-26 11:00:00	True	2022-03-30 11:00:00	1	
1	2022-09-01 13:00:00	False	NaT	1	
2	2023-02-22 07:00:00	True	2023-02-24 07:00:00	1	
3	2023-01-14 08:00:00	True	2023-01-24 08:00:00	1	
4	2023-03-25 08:00:00	True	2023-03-29 08:00:00	1	
5	2023-05-26 07:00:00	True	2023-05-27 07:00:00	1	
6	2023-06-22 11:00:00	True	2023-06-26 11:00:00	1	
7	2023-06-20 15:00:00	False	NaT	1	

	checked_bags	trip_airline	destination	user_id	home_city
0	1	United Airlines	los angeles	49728	long beach
1	1	Delta Air Lines	los angeles	245954	santa ana
2	0	Delta Air Lines	los angeles	507931	long beach
3	2	Volaris	los angeles	510094	long beach
4	0	United Airlines	los angeles	624064	long beach
5	1	Delta Air Lines	los angeles	721200	long beach
6	1	American Airlines	new york	789756	newark
7	1	Southwest Airlines	milwaukee	920115	chicago

Interesting. I guess, \$5.35 would be cheaper than a taxi in these cases. But really?

Back to the Budget-Conscious Travelers. We need to filter the fares a bit to exclude some outliers. We will use the 1st and 99th percentiles

```
In [31]: # Calculate the 1st and 99th percentiles
lower_percentile = df_avg_base_fare['avg_base_fare'].quantile(0.01)
upper_percentile = df_avg_base_fare['avg_base_fare'].quantile(0.99)

# Filter out values below the 1st percentile and above the 99th percentile
df_filtered_avg_base_fare = df_avg_base_fare[(df_avg_base_fare['avg_base_fare'] >= lower_percentile) &
                                              (df_avg_base_fare['avg_base_fare'] <= upper_percentile)]

# Display the filtered DataFrame statistics
print(df_filtered_avg_base_fare['avg_base_fare'].describe())
```

```
count    5100.000000
mean      491.411119
std       395.947058
min        62.410000
25%       286.955833
50%       392.895000
75%       559.405000
max       3449.390000
Name: avg_base_fare, dtype: float64
```

\$62.41 is still low, but within reason for some of the budget airlines.

Now a ratio of fare to distance. This can help determine which fares are lower or higher than average